

Java Collections Framework - Detailed Guide

The Java Collections Framework (JCF) provides a set of interfaces and classes to store, manipulate, and retrieve data efficiently. The diagram illustrates the hierarchy of key interfaces and classes in the framework. Below, we describe each structure, its time complexity, and when to use it

| Structure / Class | Random access | Add (end) | Add (middle) | Remove (by index) | Remove (value) | Contains / Search | Iteration |
|-------------------|-------------------------|-------------------------------|--------------|-------------------|----------------|-------------------|------------------------|
| ArrayList | O(1) | Amortized O(1) | O(n) | O(n) | O(n) | O(n) | O(n) |
| LinkedList | O(n) | O(1) (at ends) | O(n) | O(n) | O(n) | O(n) | O(n) |
| Vector | O(1) | Amortized O(1) (synchronized) | O(n) | O(n) | O(n) | O(n) | O(n) |
| Stack (legacy) | O(1) | O(1) | — | O(1) | — | — | O(n) |
| ArrayDeque | no random access (O(n)) | O(1) | O(n) | O(n) | O(n) | O(n) | |
| PriorityQueue | — | O(log n) | — | O(n) | O(n) | O(n) | O(n) |
| HashSet | — | O(1) avg | — | — | O(1) avg | O(1) avg | O(n) |
| LinkedHashSet | — | O(1) avg | — | — | O(1) avg | O(1) avg | O(n) (insertion order) |
| TreeSet | — | O(log n) | — | — | O(log n) | O(log n) | O(n) (sorted order) |

Detailed Descriptions & When to Use

Iterable & Collection (interfaces)

- **Purpose:** `Iterable<T>` is the base interface that allows objects to be iterated with `for-each`. `Collection<E>` extends `Iterable` and is the root for groups of objects (`List`, `Set`, `Queue`).
- **When to use:** Use these interface types for API signatures so you accept multiple implementations.

List (interface)

- **Properties:** Ordered collection, allows duplicates, positional access via index.
- **Common implementations:** `ArrayList`, `LinkedList`, `Vector`.
- **Use when:** You need an ordered sequence with positional access or duplicates allowed.

ArrayList (class)

- **Internal structure:** Resizable array.
- **Performance:** `get(i)` $O(1)$; add at end amortized $O(1)$; insert/remove at middle $O(n)$.
- **Memory:** Compact contiguous array; requires resizing growth (usually by $\sim 1.5\text{--}2\times$).
- **When to use:** Default go-to list. Use when you need fast random access and mostly append or iterate.
- **Caveat:** Inserting/removing from front or middle is expensive.

LinkedList (class)

- **Internal structure:** Doubly-linked list; implements `List` and `Deque`.
- **Performance:** Fast inserts/removes at ends $O(1)$; random access $O(n)$.
- **When to use:** Use when you must frequently insert/remove from both ends or when you have frequent removals via iterators. Rarely faster than `ArrayList` in practice due to cache inefficiency.
- **Caveat:** Higher memory per element (node overhead).

Vector & Stack (legacy)

- **Vector:** Synchronized, similar API to `ArrayList`. Avoid unless you require legacy synchronized behavior.
- **Stack:** Extends `Vector`. Historically used for LIFO, but prefer `ArrayDeque` or `Deque` interface as a modern stack.

Queue & Deque (interfaces)

- **Queue:** FIFO semantics; implementations include `LinkedList`, `ArrayDeque`, `PriorityQueue`.
- **Deque:** Double-ended queue — add/remove from both ends. Implementations: `ArrayDeque`, `LinkedList`.

ArrayDeque

- **Internal structure:** Circular resizable array.
- **Performance:** $O(1)$ add/remove at both ends; faster than LinkedList in practice for deque operations.
- **When to use:** Use as a queue or stack (push/pop) when you need high performance, and null elements are not required (it does **not** allow null).
- **Preferred over:** LinkedList for queue/stack roles in most cases.

PriorityQueue

- **Internal structure:** Binary heap.
- **Performance:** add $O(\log n)$, poll $O(\log n)$, peek $O(1)$.
- **Ordering:** Natural ordering or custom Comparator. Not sorted iteration — only head is guaranteed minimal/maximal.
- **When to use:** Scheduling tasks, Dijkstra-like algorithms, event simulation — anytime you need to always access the smallest/largest element quickly.
- **Caveat:** Does not permit null elements.

Set (interface)

- **Properties:** No duplicate elements. No guaranteed order (depends on implementation).
- **Common implementations:** HashSet, LinkedHashSet, TreeSet.

HashSet

- **Internal structure:** Backed by a HashMap.
- **Performance:** Average $O(1)$ for add/remove/contains.
- **When to use:** Fast membership testing, uniqueness enforcement without order requirement.
- **Nulls:** Allows one null.

LinkedHashSet

- **Properties:** Hash-based but maintains insertion (or access) order via linked list.
- **When to use:** When deterministic iteration order is required in addition to $O(1)$ operations.

TreeSet

- **Internal structure:** Red-black tree (TreeMap backing).
- **Performance:** $O(\log n)$ for add/remove/contains.
- **Ordering:** Sorted order (natural or via comparator). Supports range operations (subSet, headSet, tailSet).
- **When to use:** Need sorted unique elements, navigable set operations, or range queries.
- **Nulls:** Usually disallowed unless comparator handles null.

Table 1 — Memory & Practical Performance Notes

| Comparison | Memory Overhead | Speed Notes | Recommendation |
|---------------------------------|--|--|--|
| ArrayList vs LinkedList | ArrayList: low overhead (contiguous array). LinkedList: high overhead (node object + 2 pointers per element). | ArrayList faster for iteration & random access; LinkedList only better for frequent insert/remove at ends. | Default to ArrayList unless you truly need frequent end insertions/removals. |
| ArrayDeque vs LinkedList | ArrayDeque: low overhead (circular array). LinkedList: high overhead. | ArrayDeque typically faster for queue/stack ops due to better cache locality. | Prefer ArrayDeque for most double-ended queue needs. |
| HashSet vs TreeSet | HashSet: lower overhead, based on hash table. TreeSet: more overhead (tree nodes). | HashSet $O(1)$ avg ops; TreeSet $O(\log n)$ but provides sorted order & range queries. | Use HashSet for raw speed, TreeSet for ordering. |
| Null Elements | HashSet allows 1 null; ArrayList allows multiple nulls; PriorityQueue, ArrayDeque, TreeSet usually disallow. | — | Always check null element policy before choosing. |

Table 2 — Concurrency & Thread Safety

| Approach / Class | Thread Safety | Performance Impact | When to Use |
|---|--|--|---|
| ArrayList, HashSet, etc. | Not thread-safe | Highest performance in single-threaded contexts | Use in single-threaded or externally synchronized code |
| Collections.synchronizedList / synchronizedSet | Thread-safe (via synchronized wrapper) | Adds locking overhead | For small-scale multithreading where simple sync is enough |
| Vector / Stack (legacy) | Thread-safe (synchronized methods) | Higher overhead, outdated | Maintain legacy code requiring these types; avoid in new code |
| ConcurrentHashMap | Thread-safe (concurrent locks) | High scalability for maps | For high-performance concurrent map usage |
| CopyOnWriteArrayList | Thread-safe (copy-on-write) | Good for read-heavy, bad for write-heavy workloads | Use when reads dominate and modifications are rare |
| ConcurrentLinkedQueue | Thread-safe (non-blocking) | Scales well for concurrent queues | Use in producer-consumer models |